

# Logistic Regression & Tree Classification

Group 4 - Cullen McNamee, Josep Nueno, and WanQi Tay

8/25/2018

```
suppressWarnings(library(dplyr))
```

```
##  
## Attaching package: 'dplyr'  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
suppressWarnings(library(gains))  
suppressWarnings(library(AUC))
```

```
## AUC 0.3.0
```

```
## Type AUCNews() to see the change log and ?AUC to get an overview.
```

```
suppressWarnings(library(rpart))  
suppressWarnings(library(rpart.plot))
```

## Data Preparation

```
melbourne <- read.csv('melbourne_edited.csv')
```

```
#Identified categorical variables by sorting by the number of unique values.
```

```
uniqueCounts <- apply(melbourne, 2, unique)  
sort(unlist(lapply(uniqueCounts, length)))
```

```
##           Type      Regionname      Method      Rooms      Bathroom  
##           3           8           9           12           12  
##      Bedroom2      Car      CouncilArea      Date      YearBuilt  
##           16           16           33           78           161  
##      Postcode      Distance      Propertycount      Suburb      SellerG  
##           211           215           342           350           388  
##      BuildingArea      Landsize      Price      Longitude      Latitude  
##           741           1685           2872           5588           13403  
##           Address  
##           34006
```

```
#Exclude irrelevant or distinct fields (i.e. address and seller)
```

```
fieldsToExclude <- c('Address', 'SellerG', 'Postcode', 'Longitude',  
                    'Latitude', 'Date', 'Suburb', 'CouncilArea')
```

```
melbourne <- melbourne[,!colnames(melbourne) %in% fieldsToExclude]
```

```

#Remove NAs in Price, since this will be our Response variable.

melbourne <- melbourne %>% filter(Price != 'NA')

#Remove variables with NA in other important fields.

melbourne <- melbourne %>% filter(is.na(Bedroom2) == FALSE,
                                is.na(Bathroom) == FALSE,
                                is.na(Car) == FALSE,
                                is.na(Landsize) == FALSE,
                                is.na(BuildingArea) == FALSE,
                                is.na(YearBuilt) == FALSE)

#This leaves us with 8895 observations remaining

nrow(melbourne)

## [1] 8895

#Identify top quartile price, for use in logistic regression.
#Note: Top quartile here is an arbitrary cutoff point for classification.

summary(melbourne$Price)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 131000  640500   900000 1092524 1345000 9000000

#Code top-quartile prices as 1's and others as 0s.

melbourne$Price <- 1*(melbourne$Price > 1345000)

```

## Logistic Regression

```

#Split data set into numeric and categorical variables.

numericVariables <- c('Landsize', 'Rooms', 'Bathroom', 'Bedroom2',
                     'Car', 'Distance', 'Propertycount', 'BuildingArea', 'YearBuilt')

categoricalVariables <- c('Price', 'Type', 'Method', 'Regionname')

#Note: Price was transforming to a categorical variable: {top-quartile, not top-quartile}

melbourne[,categoricalVariables] <- lapply(melbourne[,categoricalVariables], factor)
melbourne[,numericVariables] <- apply(melbourne[,numericVariables], 2, as.numeric)

#Randomly split data set into training and holdout samples.

set.seed(10101)
randomIndex <- sample(1:nrow(melbourne), size = 0.632*nrow(melbourne))
melbourneTrain <- melbourne[randomIndex, ]
melbourneHoldout <- melbourne[-randomIndex, ]

#Run logistic regression on the training sample, using all variables.

logisticAll <- glm(Price ~ ., data = melbourneTrain, family=binomial(link=logit))

```

```

#Use step to minimize the AIC of the model, using direction both.
stepAICModel <- step(logisticAll, direction = 'both', trace = FALSE)

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

print(stepAICModel$call)

## glm(formula = Price ~ Rooms + Type + Distance + Bathroom + Car +
##      Landsize + BuildingArea + YearBuilt + Regionname + Propertycount,
##      family = binomial(link = logit), data = melbourneTrain)
logisticStepped <- stepAICModel
summary(stepAICModel)

##
## Call:
## glm(formula = Price ~ Rooms + Type + Distance + Bathroom + Car +
##      Landsize + BuildingArea + YearBuilt + Regionname + Propertycount,
##      family = binomial(link = logit), data = melbourneTrain)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0193  -0.3722  -0.1158   0.0000   3.7617
##
## Coefficients:
##              Estimate Std. Error z value
## (Intercept)    2.299e+01  3.078e+00   7.470
## Rooms          6.901e-01  7.976e-02   8.653
## Typet         -1.553e+00  2.104e-01  -7.384
## Typeu         -4.357e+00  3.895e-01 -11.187
## Distance      -3.187e-01  1.676e-02 -19.017
## Bathroom       6.155e-01  8.959e-02   6.870
## Car           2.266e-01  5.320e-02   4.260

```

```
## Landsize 1.895e-04 4.148e-05 4.569
## BuildingArea 8.419e-03 8.457e-04 9.954
## YearBuilt -1.297e-02 1.614e-03 -8.036
## RegionnameEastern Victoria -1.085e+01 6.030e+02 -0.018
## RegionnameNorthern Metropolitan -2.109e+00 1.898e-01 -11.117
## RegionnameNorthern Victoria -1.276e+01 4.499e+02 -0.028
## RegionnameSouth-Eastern Metropolitan 1.006e+00 3.107e-01 3.237
## RegionnameSouthern Metropolitan 1.068e+00 1.563e-01 6.831
## RegionnameWestern Metropolitan -2.290e+00 1.814e-01 -12.626
## RegionnameWestern Victoria -1.583e+01 7.652e+02 -0.021
## Propertycount 2.955e-05 1.309e-05 2.258
## Pr(>|z|)
## (Intercept) 8.05e-14 ***
## Rooms < 2e-16 ***
## Typet 1.53e-13 ***
## Typeu < 2e-16 ***
## Distance < 2e-16 ***
## Bathroom 6.41e-12 ***
## Car 2.04e-05 ***
## Landsize 4.90e-06 ***
## BuildingArea < 2e-16 ***
## YearBuilt 9.27e-16 ***
## RegionnameEastern Victoria 0.98564
## RegionnameNorthern Metropolitan < 2e-16 ***
## RegionnameNorthern Victoria 0.97737
## RegionnameSouth-Eastern Metropolitan 0.00121 **
## RegionnameSouthern Metropolitan 8.45e-12 ***
## RegionnameWestern Metropolitan < 2e-16 ***
## RegionnameWestern Victoria 0.98350
## Propertycount 0.02395 *
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6270.2 on 5620 degrees of freedom
## Residual deviance: 2833.5 on 5603 degrees of freedom
## AIC: 2869.5
##
## Number of Fisher Scoring iterations: 16
```

*#Using probability cutoff of 0.5, predict 1s or 0s for Price.*

```
fittedValues <- logisticStepped$fitted.values
fittedValues[fittedValues>=0.5]=1
fittedValues[fittedValues<0.5]=0
```

### Confusion Matrix of Train

```
table(melbourneTrain$Price,fittedValues)
```

```
## fittedValues
## 0 1
## 0 4007 232
## 1 382 1000
```

```
round(prop.table(table(melbourneTrain$Price,fittedValues),1),2)
```

```
##      fittedValues
##           0      1
##    0 0.95 0.05
##    1 0.28 0.72
```

```
#Run predictions on holdout sample.
```

```
predictions.holdout=predict(logisticStepped, newdata=melbourneHoldout[, -3], type="response")
predictedValues <- predictions.holdout
predictedValues[predictedValues>=0.5]=1
predictedValues[predictedValues<0.5]=0
```

### Confusion Matrix of Holdout

```
table(melbourneHoldout$Price,predictedValues)
```

```
##      predictedValues
##           0      1
##    0 2295  140
##    1  235  604
```

```
round(prop.table(table(melbourneHoldout$Price,predictedValues),1),2)
```

```
##      predictedValues
##           0      1
##    0 0.94 0.06
##    1 0.28 0.72
```

### Gains Table


```
gains(as.numeric(melbourneHoldout$Price)-1,predictions.holdout,10)
```

##	Depth				Cume	Cume Pct			Mean
##	of		Cume	Mean	Mean	of Total	Lift	Cume	Model
##	File	N	N	Resp	Resp	Resp	Index	Lift	Score
##	10	327	327	0.94	0.94	36.8%	369	369	0.96
##	20	327	654	0.75	0.85	65.9%	291	330	0.75
##	30	328	982	0.45	0.71	83.7%	177	279	0.44
##	40	327	1309	0.26	0.60	93.7%	100	234	0.23
##	50	328	1637	0.10	0.50	97.5%	38	195	0.11
##	60	327	1964	0.03	0.42	98.8%	13	165	0.05
##	70	327	2291	0.02	0.36	99.4%	6	142	0.02
##	80	328	2619	0.01	0.32	99.9%	5	125	0.01
##	90	327	2946	0.00	0.28	99.9%	0	111	0.00
##	100	328	3274	0.00	0.26	100.0%	1	100	0.00

### Gains Plot

```
plot(gains(as.numeric(melbourneHoldout$Price)-1,
           predictions.holdout,10), ylim = c(0, 1.3),
      xlim = c(0, 100))
```


**Gains Table Plot**



**AUC Plot**

```
plot(roc(predictions.holdout, melbourneHoldout$Price), main = "AUROC")
```

**AUROC**




## Tree Classification Model

```
#Fit tree model using rpart.
#Minimum split of 30.

classificationTrain <-
  rpart(Price ~.,data = melbourneTrain,
        control=rpart.control(cp=0,minsplit=30,xval=10, maxsurrogate=0))
printcp(classificationTrain)

##
## Classification tree:
## rpart(formula = Price ~ ., data = melbourneTrain, control = rpart.control(cp = 0,
##   minsplit = 30, xval = 10, maxsurrogate = 0))
##
## Variables actually used in tree construction:
## [1] Bathroom      Bedroom2      BuildingArea  Car           Distance
## [6] Landsize      Propertycount Regionname    Type          YearBuilt
##
## Root node error: 1382/5621 = 0.24586
##
## n= 5621
##
##      CP nsplit rel error  xerror    xstd
## 1  0.17329957      0  1.00000 1.00000 0.023360
## 2  0.07525326      2  0.65340 0.65412 0.019930
## 3  0.02484322      3  0.57815 0.58466 0.019033
## 4  0.01447178      6  0.50362 0.51954 0.018108
## 5  0.01338640      8  0.47467 0.51447 0.018033
## 6  0.01266281     10  0.44790 0.51158 0.017989
## 7  0.01085384     12  0.42258 0.48770 0.017623
## 8  0.01013025     13  0.41172 0.47250 0.017383
## 9  0.00361795     14  0.40159 0.44645 0.016958
## 10 0.00325615     15  0.39797 0.43488 0.016764
## 11 0.00289436     21  0.37410 0.43849 0.016825
## 12 0.00180897     25  0.36252 0.44211 0.016886
## 13 0.00168837     31  0.35166 0.43705 0.016801
## 14 0.00144718     34  0.34660 0.43271 0.016727
## 15 0.00108538     42  0.33213 0.42909 0.016665
## 16 0.00072359     48  0.32562 0.43632 0.016788
## 17 0.00036179     51  0.32344 0.45658 0.017126
## 18 0.00024120     53  0.32272 0.46093 0.017197
## 19 0.00000000     56  0.32200 0.46454 0.017255

plotcp(classificationTrain, minline = TRUE)
```



*#Find minimum error of CP.*

```
cpResults <- printcp(classificationTrain)
```

```
##
## Classification tree:
## rpart(formula = Price ~ ., data = melbourneTrain, control = rpart.control(cp = 0,
##   minsplit = 30, xval = 10, maxsurrogate = 0))
##
## Variables actually used in tree construction:
## [1] Bathroom      Bedroom2      BuildingArea  Car           Distance
## [6] Landsize      Propertycount Regionname    Type          YearBuilt
##
## Root node error: 1382/5621 = 0.24586
##
## n= 5621
##
##      CP nsplit rel error  xerror   xstd
## 1  0.17329957      0  1.00000 1.00000 0.023360
## 2  0.07525326      2  0.65340 0.65412 0.019930
## 3  0.02484322      3  0.57815 0.58466 0.019033
## 4  0.01447178      6  0.50362 0.51954 0.018108
## 5  0.01338640      8  0.47467 0.51447 0.018033
## 6  0.01266281     10  0.44790 0.51158 0.017989
## 7  0.01085384     12  0.42258 0.48770 0.017623
## 8  0.01013025     13  0.41172 0.47250 0.017383
## 9  0.00361795     14  0.40159 0.44645 0.016958
## 10 0.00325615     15  0.39797 0.43488 0.016764
```



```
## 11 0.00289436      21    0.37410 0.43849 0.016825
## 12 0.00180897      25    0.36252 0.44211 0.016886
## 13 0.00168837      31    0.35166 0.43705 0.016801
## 14 0.00144718      34    0.34660 0.43271 0.016727
## 15 0.00108538      42    0.33213 0.42909 0.016665
## 16 0.00072359      48    0.32562 0.43632 0.016788
## 17 0.00036179      51    0.32344 0.45658 0.017126
## 18 0.00024120      53    0.32272 0.46093 0.017197
## 19 0.00000000      56    0.32200 0.46454 0.017255
```

```
minErrorCP <- cpResults[as.numeric(which.min(cpResults[,4])),1]
print(minErrorCP)
```

```
## [1] 0.001085384
```

```
#Prune initial tree using above cp.
```

```
classificationTrainPruned <- rpart(Price ~., data = melbourneTrain,
                                   control=rpart.control(cp=minErrorCP,minsplit=30,
                                                         xval=10, maxsurrogate=0))
printcp(classificationTrainPruned)
```

```
##
```

```
## Classification tree:
```

```
## rpart(formula = Price ~ ., data = melbourneTrain, control = rpart.control(cp = minErrorCP,
##   minsplit = 30, xval = 10, maxsurrogate = 0))
##
```

```
## Variables actually used in tree construction:
```

```
## [1] Bathroom      Bedroom2      BuildingArea  Car           Distance
## [6] Landsize       Propertycount Regionname    Type          YearBuilt
##
```

```
## Root node error: 1382/5621 = 0.24586
```

```
##
```


```
## n= 5621
```

```
##
```

```
##          CP nsplit rel error  xerror    xstd
## 1  0.1732996      0  1.00000 1.00000 0.023360
## 2  0.0752533      2  0.65340 0.65412 0.019930
## 3  0.0248432      3  0.57815 0.59479 0.019169
## 4  0.0144718      6  0.50362 0.53039 0.018268
## 5  0.0133864      8  0.47467 0.52605 0.018205
## 6  0.0126628     10  0.44790 0.51881 0.018098
## 7  0.0108538     12  0.42258 0.49566 0.017747
## 8  0.0101302     13  0.41172 0.48770 0.017623
## 9  0.0036179     14  0.40159 0.44935 0.017007
## 10 0.0032562     15  0.39797 0.44428 0.016922
## 11 0.0028944     21  0.37410 0.44645 0.016958
## 12 0.0018090     25  0.36252 0.44067 0.016862
## 13 0.0016884     31  0.35166 0.45080 0.017030
## 14 0.0014472     34  0.34660 0.45876 0.017161
## 15 0.0010854     42  0.33213 0.46237 0.017220
## 16 0.0010854     48  0.32562 0.46744 0.017302
```

**CP Plot**

```
plotcp(classificationTrainPruned, minline = TRUE)
```



Confusion Matrix of Un-Pruned Tree (Train)

```
table(melbourneTrain$Price, predict(classificationTrain, type="class"))
```

```
##
##      0      1
## 0 4073  166
## 1  279 1103
```

```
round(prop.table(table(melbourneTrain$Price, predict(classificationTrain,
                                                         type="class"))), 1), 2)
```

```
##
##      0      1
## 0 0.96 0.04
## 1 0.20 0.80
```

Confusion Matrix of Pruned Tree (Train)

```
table(melbourneTrain$Price, predict(classificationTrainPruned, type="class"))
```

```
##
##      0      1
## 0 4079  160
## 1  290 1092
```

```
round(prop.table(table(melbourneTrain$Price, predict(classificationTrainPruned,
                                                         type="class"))), 1), 2)
```

```
##
##      0      1
##  0 0.96 0.04
##  1 0.21 0.79
```

### Confusion Matrix of Pruned Tree (Holdout)

```
table(melbourneHoldout[,3],
      predict(classificationTrainPruned,newdata=melbourneHoldout[, -3],type="class"))
```

```
##
##      0      1
##  0 2268  167
##  1  209  630
```

```
round(prop.table(table(melbourneHoldout[,3],
                      predict(classificationTrainPruned,newdata=melbourneHoldout[, -3],
                              type="class")),1),2)
```

```
##
##      0      1
##  0 0.93 0.07
##  1 0.25 0.75
```

### Tree Plot

*#For this analysis CouncilArea is excluded.*  
*#Min split of 500 is used.*

```
melbourneTrainWithoutCouncil <- melbourneTrain[, !colnames(melbourneTrain) %in% 'CouncilArea']

classificationTrainPruned_ExMelbourne <-
  rpart(Price ~., data = melbourneTrainWithoutCouncil,
        control=rpart.control(cp=minErrorCP, minsplit=500, xval=10, maxsurrogate=0))

rpart.plot(classificationTrainPruned_ExMelbourne, type = 5)
```

